Leveraging TF-IDF and BERT for Hate Speech Detection on Twitter

Ashanti Boone
Department of Computer Science
Bowie State University

Abstract

This research explores a methodology for detecting hate speech on social media using Natural Language Processing (NLP) techniques. Utilizing publicly available Twitter datasets from Kaggle, the approach involves preprocessing textual data, extracting features, and employing dictionaries of hateful words for classification. Key NLP techniques such as tokenization, lemmatization, and word frequency analysis are integrated to enhance detection accuracy. Aimed at addressing the pressing issue of cyberbullying, particularly among vulnerable groups like children, the research seeks to enable real-time moderation of harmful content and foster safer digital spaces. The system's effectiveness and scalability are demonstrated through its implementation and evaluation on the datasets, highlighting its potential for real-world applications on social media platforms.

Introduction

Definition and Concern: Hate speech involves communication that *demeans, insults, or incites violence* against individuals or groups based on attributes like race, religion, ethnicity, gender, or sexual orientation. In the digital era, its impact is amplified by social media platforms like Twitter, enabling harmful content to spread rapidly.

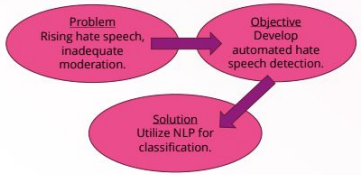
Challenges of Moderation:
•Hate speech often *goes viral*, increasing its reach and potential to incite real-world violence.
•Its nuanced nature, including the use of *slang, sarcasm, and context dependency*, complicates detection.
•Vulnerable populations, such as *youth exposed to cyberbullying*, are particularly affected.

Limitations of Manual Methods: Traditional content moderation relies on *manual reviews*, which are inefficient given the vast scale of online content.

Automated Solutions:
•Researchers leverage *Natural Language Processing (NLP)* and machine learning for real-time hate speech detection.
•Techniques like *TF-IDF*, word embeddings, and *deep learning models like BERT* are employed to enhance accuracy.

Impact: A well-designed automated system fosters *safer online environments* and strengthens content moderation efforts to combat the spread of hate speech effectively.

Problem Statement



Methodology

Techniques and Approach:
•Utilizes *Natural Language Processing (NLP)* and machine learning for tweet hate speech detection.
•Textual data was *preprocessed* and converted into numerical representations using **TF-IDF (Term Frequency-Inverse Document Frequency)** to highlight the importance of unique terms.

Feature Selection: A vocabulary of the *178 most significant terms* was chosen to balance *computational efficiency* and informativeness.

Model Training: Initial classifiers included *logistic regression* and *support vector machines*.

BERT embeddings were incorporated to improve performance by capturing *contextual relationships* between words.

Hybrid Approach: Combined *term-based importance* (TF-IDF) with *contextual comprehension* (BERT).

Evaluation and Optimization: Model performance was measured using accuracy, F1-score, and recall metrics. *Parameter tuning* and iterative refinements guided the optimization process.

Dataset

This research utilized two primary datasets from Kaggle, provided by *hate_speech_offensive*. The first dataset contained 10,000 tweets with metadata such as usernames, interactions, and timestamps, while the second included 24,783 tweets labeled by three annotators as *hate speech, offensive language*, or *neither*. These labels were crucial for supervised learning and model evaluation. Additionally, a hate speech dictionary sourced from GitHub by user *t-davidson* was incorporated. This dictionary comprised 176 hate speech terms, as a foundation for feature extraction and keyword analysis in detecting harmful content.

Challenges

Context Misclassification: Tweets with terms from the hate speech dictionary were flagged as hate speech even when the context was not harmful.
- Example: Phrases like *"white boy"* were incorrectly categorized due to **keyword-matching limitations** that overlooked contextual nuances.

Dataset Bias:
•**Subjective judgments** by annotators led to inconsistencies in labeling tweets as hate speech, offensive, or neutral.
•The variability of hate speech definitions influenced data reliability and model performance.

Imbalanced Dataset: Non-hateful content significantly *outnumbered* hateful tweets. This impacted model training, skewing performance metrics and limiting the ability to generalize detection capabilities.

Key Takeaways:
•Highlighted the need for **context-aware approaches** to address misclassification challenges.
•Called for **refined data processing** and **model enhancements** to improve the reliability and fairness of hate speech detection systems.

Fly white person consumer list.

Example of a false positive.

Results

Round 1 (Initial Experimentation):
•**Incomplete labeling** for hate speech, offensive, and neutral categories caused skewed results.
•**Inflated Accuracy**:
- Hate speech accuracy reached **97%**, while non-hate speech accuracy was only **41%**, yielding an overall accuracy of **77%**.

Post-Correction Rounds:
•Adjusted class weights and improved labeling led to **stabilized overall accuracy**:

Class Performance:
•**Offensive Language**:
- Precision and recall consistently exceeded **93%**, demonstrating high reliability.

•**"Neither" Category**:
- Precision and recall ranged from **85% to 89%**, effectively identifying non-harmful content.

•**Hate Speech**:
- Precision ranged between **46.67% and 48.95%**, with recall peaking at **41.03%** in Round 3.
- Performance struggled due to the **dataset's limited representation of hate speech tweets**, emphasizing the challenges of detecting nuanced language.

Key Insights:
•The model excelled in detecting offensive language and neutral content.
•However, **class imbalance** and the **complexity of hate speech detection** impacted the precision and recall for this critical category..

Training Round	Overall	Hate Speech	Offensive	Neither
Pre	77%	97%	N/A	41%
1	90%	40%	95%	88%
2	90%	39%	95%	88%
3	90%	41%	95%	85%

Summary table of detection accuracy for four rounds of BERT training.





Conclusion

Initial Pre-Round Results: **Inflated Accuracy**: Achieved **97% for hate speech detection** due to improperly labeled categories, emphasizing the critical need for **accurate labeling** and **balanced datasets**.

Refined Model Performance:
•**Overall Accuracies**: Exceeded **90%** across three training rounds, demonstrating reliability for **offensive language** and **neutral tweet classification**.

Hate Speech Detection Metrics:
•**F1-Score**: Ranged from **42.88% to 44.02%**, reflecting moderate effectiveness in balancing **precision** and **recall**.
•**Precision**: Measured between **46.67% and 48.95%**, highlighting the proportion of correctly identified hate speech.
•**Recall**: Varied from **38.97% to 41.03%**, indicating the model's ability to capture all hate speech instances.
•These metrics underscore the **challenges of detecting nuanced language** in a dataset with a **limited representation** of hate speech.

Key Takeaways:
•The model is a **foundational solution** for hate speech detection rather than a definitive one.
•While it demonstrates promise, targeted improvements are necessary to enhance its performance for hate speech in **real-world applications**.

Recommendations

To enhance hate speech detection, future efforts should focus on expanding and balancing the dataset to address the limited representation of hate speech instances, ensuring the model is trained on a more diverse and comprehensive set of examples. Additionally, optimization techniques should prioritize improving precision, recall, and f1-score for hate speech classification to better capture nuanced and contextual language. Cross-domain validation, using datasets from multiple platforms, can further test the model's adaptability and robustness in detecting hate speech across varied social media environments.

Expand and Balance The Dataset

Optimize For Hate Speech Metrics

Cross Domain Validation

Acknowledgements

Research supported by Bowie State University
Collaborators: Vivek Shandilya | Department of Computer Science